



ITESO, Universidad
Jesuita de Guadalajara

✓ Deep Learning.

Proyecto 2 - GANs.

Nombre: Sergio Daniel Dueñas Godinez.

Expediente : 739300.

Profesor: Iván Reyes Amezcua.

Link Github: <https://github.com/SergioDuenass/Face-Aging>

✓ Objetivos:

- Comprender los principios fundamentales de las redes generativas antagónicas (GANs) y su aplicación en el aprendizaje profundo generativo.
- Implementar una GAN básica y explorar variantes como las GANs condicionales o las CycleGANs para tareas específicas, como la generación de imágenes, traducción de imágenes o generación de texto.
- Analizar el rendimiento y las características de los modelos generados por las GANs.

Descripción:

Deberán seleccionar un conjunto de datos adecuado para su proyecto, que puede ser de imágenes, texto o cualquier otro tipo que permita la aplicación de GANs. Implementar una GAN, como una GAN básica, una GAN condicional o una CycleGAN, dependiendo de la naturaleza del conjunto de datos y el objetivo del proyecto. El proyecto incluirá una fase de experimentación donde los estudiantes deberán entrenar, ajustar y evaluar sus modelos. Presentar sus resultados a través de un informe escrito y una presentación, discutiendo la implementación, los desafíos encontrados, el rendimiento de sus modelos y las aplicaciones potenciales de su trabajo.

Lo que yo haré será una cGAN, una GAN condicional, en donde se buscará a una cara, ponerla en rango de edad que va de 10 en 10.

Los datos son recopilados de <https://data.vision.ee.ethz.ch/cvl/rrothe/imdb-wiki/> Ya que estos datos son los que mejor se acoplarían a nuestro modelo, ya vienen listos con su label ya hechos

```
# Librerías a utilizar
import math
import os
import time
import numpy as np
import pickle as pkl
from PIL import Image
from glob import glob
from pathlib import Path
from scipy.io import loadmat
from datetime import datetime
import matplotlib.pyplot as plt
```

```
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torch.utils.data import Dataset, DataLoader
from torchvision import datasets, transforms, utils
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
%matplotlib inline
```

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
device

    device(type='cuda')
```

▼ Obtención de datos

```
!wget https://data.vision.ee.ethz.ch/cvl/rrothe/imdb-wiki/static/wiki_crop.tar
!tar -xf wiki_crop.tar

--2024-04-11 21:25:18--  https://data.vision.ee.ethz.ch/cvl/rrothe/imdb-wiki/
Resolving data.vision.ee.ethz.ch (data.vision.ee.ethz.ch)... 129.132.52.178, :
Connecting to data.vision.ee.ethz.ch (data.vision.ee.ethz.ch)|129.132.52.178|
```

```
HTTP request sent, awaiting response... 200 OK
Length: 811315200 (774M) [application/x-tar]
Saving to: 'wiki_crop.tar'
```

```
wiki_crop.tar      100%[=====>] 773.73M  15.3MB/s   in 52s
```

```
2024-04-11 21:26:11 (14.9 MB/s) - 'wiki_crop.tar' saved [811315200/811315200]
```

▼ Preprocesamiento

Esta función toma la fecha de nacimiento y el año actual como entrada, y devuelve la edad en años. La lógica adicional se encarga de manejar casos donde el cumpleaños aún no ha ocurrido en el año actual.

```
def calcular_edad(taken, dob):

    birth = datetime.fromordinal(max(int(dob) - 366, 1))

    if birth.month < 7:
        return taken - birth.year
    else:
        return taken - birth.year - 1
```

Esta función carga y prepara los datos de imágenes y edades para su posterior procesamiento en un formato útil y limpio.

```
def load_data(dataset='wiki', data_dir='./wiki_crop'):

    # Cargamos la metadata
    meta_path = Path(data_dir) / f'{dataset}.mat'
    meta = loadmat(meta_path)
    meta_data = meta[dataset][0, 0]

    # Cargamos la lista de los paths de los archivos
    full_path = meta_data['full_path'][0]
    full_path = [y for x in full_path for y in x]

    # Nacimientos
    dob = meta_data['dob'][0]

    # Año en que se tomó la foto
    photo_taken = meta_data['photo_taken'][0]

    # Calcular la edad
    age = [calcular_edad(photo_taken[i], dob[i]) for i in range(len(dob))]
```

```
# limpiamos edad > 0
clean_mapping = {pth:age for (pth, age) in zip(full_path, age) if age > 0}

# lista de el path y la edad
full_path = list(clean_mapping.keys())
age = list(clean_mapping.values())

return full_path, age
```

Esta función simplemente hace el escalamiento para poder utilizarlo, paso importante, ya que sino quedaría imágenes sin sentido, puros colores

```
# Escalamiento de -1 a 1
def scale(x, feature_range=(-1, 1)):

    min, max = feature_range
    x = x * (max - min) + min
    return x
```

Esta función convierte un tensor x a un tensor one-hot

```
# one-hot encoder
bins = [18, 29, 39, 49, 59]
def one_hot(x, bins):

    x = x.numpy()
    idxs = np.digitize(x, bins, right=True)
    idxs = idxs.reshape(-1,1)
    z = torch.zeros(len(x), len(bins)+1).scatter_(1, torch.tensor(idxs), 1)
    return z
```

Ahora clases para la importacion de datos

esta clase ImageAgeDataset permite cargar y acceder a muestras de imágenes junto con sus edades correspondientes desde un conjunto de datos, y opcionalmente aplicar transformaciones a estas muestras.

```
class ImageAgeDataset(Dataset):

    def __init__(self, dataset, data_dir, transform=None):

        .. .. .
```

```
self.data_dir = data_dir
self.full_path, self.age = load_data(dataset, data_dir)
self.transform = transform

def __len__(self):
    return len(self.age)

def __getitem__(self, idx):
    image = Image.open(os.path.join(self.data_dir, self.full_path[idx]))
    age = self.age[idx]
    sample = {'image': image, 'age': age}
    if self.transform:
        sample = self.transform(sample)
    return sample

class Resize(object):

    def __init__(self, output_size):
        assert isinstance(output_size, (int, tuple))
        self.output_size = output_size

    def __call__(self, sample):
        image, age = sample['image'], sample['age']
        image = transforms.Resize(self.output_size)(image)
        return {'image': image, 'age': age}

class ToTensor(object):

    def __call__(self, sample):
        image, age = sample['image'], sample['age']
        image = transforms.ToTensor()(image)

        if image.size()[0] == 1:
            image = image.expand(3,-1,-1)
        return {'image': image, 'age': age}
```

Haz doble clic (o ingresa) para editar

```
dataset='wiki'
data_dir='./wiki_crop'
bins = [18, 29, 39, 49, 59]
img_size = 64
batch_size = 128

tfms = transforms.Compose([Resize((img_size, img_size)),
                           ToTensor()])

train_dataset = ImageAgeDataset(dataset, data_dir, transform=tfms)

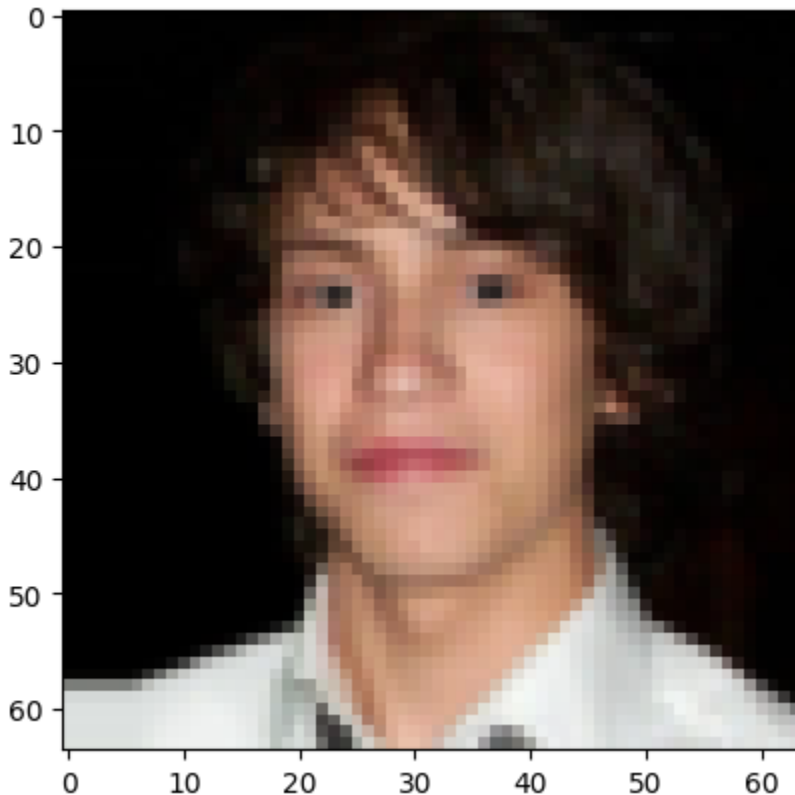
# Data loaders
```

```
# DataLoaders
train_loader = DataLoader(dataset=train_dataset, batch_size=batch_size)

for batch in train_loader:
    iter(train_loader)
```

▼ Imprimimos una muestra de los datos

```
plt.imshow(train_dataset[100]['image'].numpy().transpose(1,2,0))
<matplotlib.image.AxesImage at 0x7a6d582bffd0>
```



```
# Obtener un lote de imágenes de entrenamiento
for data in train_loader:
    images, labels = data['image'], data['age']
    break # Solo necesitamos el primer lote para este ejemplo

# Mostrar las imágenes en el lote, junto con las etiquetas correspondientes
fig = plt.figure(figsize=(25, 4))
plot_size = 20
num_cols = int(plot_size / 2) # Convertir el resultado a un entero
for idx in range(plot_size):
    ax = fig.add_subplot(2, num_cols, idx+1, xticks=[], yticks=[])
    ax.imshow(np.transpose(images[idx], (1, 2, 0)))
    # Imprimir la etiqueta correcta para cada imagen
    ax.set_title(str(labels[idx].item()))
```



✓ Ahora iniciamos los modelos

```
def conv(in_channels, out_channels, kernel_size=4, stride=2, padding=1, batch_norm=True):
    layers = []
    conv_layer = nn.Conv2d(in_channels, out_channels, kernel_size, stride, padding)
    layers.append(conv_layer)

    if batch_norm:
        bn = nn.BatchNorm2d(out_channels)
        layers.append(bn)

    return nn.Sequential(*layers)
```

✓ Discriminador

```
class Discriminator(nn.Module):

    def __init__(self, y_size, conv_dim=64):

        super(Discriminator, self).__init__()
        self.conv_dim = conv_dim
        self.y_size = y_size
        self.conv1 = conv(3, conv_dim, 4, batch_norm=False)
        self.conv2 = conv(conv_dim+y_size, conv_dim * 2, 4)
        self.conv3 = conv(conv_dim*2, conv_dim*4, 4)
        self.conv4 = conv(conv_dim*4, conv_dim*8, 4)
        self.conv5 = conv(conv_dim*8, 1, 4, 1, 0, batch_norm=False)

    def forward(self, x, y):
```

```
x = F.relu(self.conv1(x))
y = y.view(-1,y.size()[-1],1,1)
y = y.expand(-1,-1,x.size()[-2], x.size()[-1])
x = torch.cat([x, y], 1)
x = F.relu(self.conv2(x))
x = F.relu(self.conv3(x))
x = F.relu(self.conv4(x))
x = self.conv5(x)

return x
```

▼ Generador

```
def deconv(in_channels, out_channels, kernel_size=4, stride=2, padding=1, batch_n

layers = []
t_conv = nn.ConvTranspose2d(in_channels, out_channels, kernel_size, stride, p
layers.append(t_conv)

if batch_norm:
    layers.append(nn.BatchNorm2d(out_channels))

return nn.Sequential(*layers)

class Generator(nn.Module):

    def __init__(self, z_size, y_size, conv_dim=64):
        super(Generator, self).__init__()

        self.conv_dim = conv_dim

        self.t_conv1 = deconv(z_size+y_size, conv_dim*8, 4, 1, 0)
        self.t_conv2 = deconv(conv_dim*8, conv_dim*4, 4)
        self.t_conv3 = deconv(conv_dim*4, conv_dim*2, 4)
        self.t_conv4 = deconv(conv_dim*2, conv_dim, 4)
        self.t_conv5 = deconv(conv_dim, 3, 4, batch_norm=False)

    def forward(self, z, y):

        x = torch.cat([z, y], dim=1)
        x = x.view(-1, x.size()[-1], 1, 1)
        x = F.relu(self.t_conv1(x))
        x = F.relu(self.t_conv2(x))
        x = F.relu(self.t_conv3(x))
        x = F.relu(self.t_conv4(x))
        x = self.t_conv5(x)
        x = torch.tanh(x)
```



```
return x
```

▼ Lo combinamos

```
# Hiperparametros
```

```
conv_dim = 64
```

```
z_size = 100
```

```
y_size = 6 # clases
```

```
D = Discriminator(y_size, conv_dim)
```

```
G = Generator(z_size, y_size, conv_dim)
```

```
print(D)
```

```
print()
```

```
print(G)
```

```
Discriminator(
  (conv1): Sequential(
    (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=
  )
  (conv2): Sequential(
    (0): Conv2d(70, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), b
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running
  )
  (conv3): Sequential(
    (0): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), l
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running
  )
  (conv4): Sequential(
    (0): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), l
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running
  )
  (conv5): Sequential(
    (0): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
  )
)
```

```
Generator(
  (t_conv1): Sequential(
    (0): ConvTranspose2d(106, 512, kernel_size=(4, 4), stride=(1, 1), bias=Fa
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running
  )
  (t_conv2): Sequential(
    (0): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding:
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running
  )
  (t_conv3): Sequential(
    (0): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding:
    (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running
  )
  (t_conv4): Sequential(
    (0): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running:
```

```

    )
    (t_conv5): Sequential(
      (0): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1))
    )
  )
)

```

▼ Calculamos la perdida

```

def real_loss(D_out, smooth=False):
    batch_size = D_out.size(0)

    if smooth:
        labels = torch.ones(batch_size)*0.9
    else:
        labels = torch.ones(batch_size)

    labels = labels.to(device)
    criterion = nn.BCEWithLogitsLoss()

    loss = criterion(D_out.squeeze(), labels)
    return loss

def fake_loss(D_out):
    batch_size = D_out.size(0)
    labels = torch.zeros(batch_size) #
    labels = labels.to(device)
    criterion = nn.BCEWithLogitsLoss()

    loss = criterion(D_out.squeeze(), labels)
    return loss

```

▼ Optimizadores

```

# params
lr = 0.0002
beta1=0.5
beta2=0.999

# optimizadores para el generador y discirminador
d_optimizer = optim.Adam(D.parameters(), lr, [beta1, beta2])
g_optimizer = optim.Adam(G.parameters(), lr, [beta1, beta2])

```

▼ Entrenamiento

```

def checkpoint(G, D, epoch, model, root_dir):
    target_dir = f'{root_dir}/{model}'
    os.makedirs(target_dir, exist_ok=True)
    G_path = os.path.join(target_dir, f'G_{epoch}.pkl')
    D_path = os.path.join(target_dir, f'D_{epoch}.pkl')
    torch.save(G.state_dict(), G_path)
    torch.save(D.state_dict(), D_path)

def oh_to_class(fixed_y):
    age_map = {0:'0-18',1:'19-29',2:'30-39',3:'40-49',4:'50-59',5:'60+'}
    if torch.cuda.is_available():
        fixed_y = fixed_y.cpu()
    fixed_y_idxs = fixed_y.numpy().nonzero()[1]
    fixed_y_ages = [age_map[idx] for idx in fixed_y_idxs]

    return fixed_y_ages

def save_samples_ages(samples, fixed_y, model, root_dir):
    fixed_y_ages = oh_to_class(fixed_y)
    samples_ages = {'samples': samples, 'ages': fixed_y_ages}
    target_dir = f'{root_dir}/{model}'
    os.makedirs(target_dir, exist_ok=True)
    with open(f'{target_dir}/train_samples_ages.pkl', 'wb') as f:
        pkl.dump(samples_ages, f)

%%time

root_dir = '/content/Age-cGAN'
model = 'GAN_1'
os.makedirs(root_dir, exist_ok=True)

# Usamos gpu
G.to(device)
D.to(device)

import pickle as pkl

num_epochs = 50

samples = []
losses = []

print_every = 300

sample_size=16
fixed_z = np.random.uniform(-1, 1, size=(sample_size, z_size))
fixed_z = torch.from_numpy(fixed_z).float()
fixed_y = np.random.randint(len(bins), size=sample_size)
fixed_y = fixed_y.reshape(-1,1)
fixed_y = torch.zeros(sample_size, len(bins)+1).scatter_(1, torch.tensor(fixed_y)

```

```
# Entrenamos la red
for epoch in range(num_epochs):

    for batch_i, batch in enumerate(train_loader):

        batch_size = batch['image'].size(0)

        # Las reescalamos
        real_images = scale(batch['image'])

        # one-hot age
        ages = one_hot(batch['age'], bins)

# ----- Discriminador -----
    d_optimizer.zero_grad()
# Primero entrenamos con imágenes reales
    real_images = real_images.to(device)
    ages = ages.to(device)

    D_real = D(real_images, ages)
    d_real_loss = real_loss(D_real)

# Ahora con las fakes

    # Generamos las imágenes fakes
    z = np.random.uniform(-1, 1, size=(batch_size, z_size))
    z = torch.from_numpy(z).float()

    z = z.to(device)
    fake_images = G(z, ages)

    D_fake = D(fake_images, ages)
    d_fake_loss = fake_loss(D_fake)

    # metricas
    d_loss = d_real_loss + d_fake_loss
    d_loss.backward()
    d_optimizer.step()

# ----- Generador -----
    g_optimizer.zero_grad()

# Primero entrenamos con imágenes reales y labels volteados

    # Generamos las imágenes fakes
    z = np.random.uniform(-1, 1, size=(batch_size, z_size))
    z = torch.from_numpy(z).float()
    z = z.to(device)
```

Epoch [1/	50]		d_loss: 1.3925		g_loss: 1.3833
Epoch [1/	50]		d_loss: 1.5737		g_loss: 4.5406
Epoch [2/	50]		d_loss: 0.6322		g_loss: 2.7557
Epoch [2/	50]		d_loss: 0.5567		g_loss: 2.8625
Epoch [3/	50]		d_loss: 0.7253		g_loss: 4.2783
Epoch [3/	50]		d_loss: 0.4185		g_loss: 2.2750
Epoch [4/	50]		d_loss: 0.3956		g_loss: 2.6836
Epoch [4/	50]		d_loss: 0.2544		g_loss: 4.3140
Epoch [5/	50]		d_loss: 0.3985		g_loss: 3.2897
Epoch [5/	50]		d_loss: 0.3816		g_loss: 3.1451
Epoch [6/	50]		d_loss: 0.5116		g_loss: 2.2838
Epoch [6/	50]		d_loss: 0.2343		g_loss: 3.1010
Epoch [7/	50]		d_loss: 0.3401		g_loss: 4.2723
Epoch [7/	50]		d_loss: 0.1935		g_loss: 3.5154
Epoch [8/	50]		d_loss: 0.4041		g_loss: 3.5860
Epoch [8/	50]		d_loss: 0.1196		g_loss: 4.7209
Epoch [9/	50]		d_loss: 0.3283		g_loss: 4.1650
Epoch [9/	50]		d_loss: 0.0874		g_loss: 3.8748
Epoch [10/	50]		d_loss: 0.2005		g_loss: 3.2187
Epoch [10/	50]		d_loss: 0.1483		g_loss: 4.6604
Epoch [11/	50]		d_loss: 0.8275		g_loss: 8.3264
Epoch [11/	50]		d_loss: 0.2139		g_loss: 4.3889


Epoch	Step	d_loss	g_loss
Epoch	12/ 50	d_loss: 0.3192	g_loss: 5.4764
Epoch	12/ 50	d_loss: 0.1573	g_loss: 4.7697
Epoch	13/ 50	d_loss: 0.1409	g_loss: 3.7057
Epoch	13/ 50	d_loss: 0.0848	g_loss: 4.0500
Epoch	14/ 50	d_loss: 0.4799	g_loss: 5.1186
Epoch	14/ 50	d_loss: 0.2072	g_loss: 4.6562
Epoch	15/ 50	d_loss: 0.2680	g_loss: 5.0592
Epoch	15/ 50	d_loss: 0.1062	g_loss: 3.8977
Epoch	16/ 50	d_loss: 0.2526	g_loss: 5.3313
Epoch	16/ 50	d_loss: 0.1066	g_loss: 5.0206
Epoch	17/ 50	d_loss: 0.1996	g_loss: 3.5145
Epoch	17/ 50	d_loss: 0.1179	g_loss: 4.6284
Epoch	18/ 50	d_loss: 1.7350	g_loss: 1.6760
Epoch	18/ 50	d_loss: 0.0588	g_loss: 3.7427
Epoch	19/ 50	d_loss: 5.9025	g_loss: 8.2725
Epoch	19/ 50	d_loss: 0.0147	g_loss: 6.1622
Epoch	20/ 50	d_loss: 0.1558	g_loss: 4.9045
Epoch	20/ 50	d_loss: 0.0709	g_loss: 6.0949
Epoch	21/ 50	d_loss: 0.0529	g_loss: 5.4335
Epoch	21/ 50	d_loss: 0.1316	g_loss: 4.6318
Epoch	22/ 50	d_loss: 0.0416	g_loss: 5.8538
Epoch	22/ 50	d_loss: 0.0235	g_loss: 7.0229
Epoch	23/ 50	d_loss: 0.2044	g_loss: 6.0770
Epoch	23/ 50	d_loss: 1.7619	g_loss: 10.3756
Epoch	24/ 50	d_loss: 0.1311	g_loss: 6.6602
Epoch	24/ 50	d_loss: 0.0845	g_loss: 5.0989
Epoch	25/ 50	d_loss: 0.0343	g_loss: 4.7657
Epoch	25/ 50	d_loss: 0.0278	g_loss: 4.1238
Epoch	26/ 50	d_loss: 0.1694	g_loss: 5.3112
Epoch	26/ 50	d_loss: 0.0318	g_loss: 4.1340
Epoch	27/ 50	d_loss: 0.7796	g_loss: 7.8849
Epoch	27/ 50	d_loss: 0.0249	g_loss: 7.8281
Epoch	28/ 50	d_loss: 0.0505	g_loss: 6.0067
Epoch	28/ 50	d_loss: 0.2542	g_loss: 4.8042
Epoch	29/ 50	d_loss: 0.5079	g_loss: 3.7556
Epoch	29/ 50	d_loss: 0.0103	g_loss: 5.9541

▼ Loss del entrenamiento

```
fig, ax = plt.subplots()
losses = np.array(losses)
plt.plot(losses.T[0], label='Discriminator', alpha=0.5)
plt.plot(losses.T[1], label='Generator', alpha=0.5)
plt.title("Training Losses")
plt.legend()
```

<matplotlib.legend.Legend at 0x7a6d5687dff0>






▼ Calamos nuestro GAN

```
def view_samples(epoch, samples, ages):
    fig, axes = plt.subplots(figsize=(16,4), nrows=2, ncols=8, sharey=True, share
    for ax, img, age in zip(axes.flatten(), samples[epoch], ages):
        img = img.detach().cpu().numpy()
        img = np.transpose(img, (1, 2, 0))
        img = ((img +1)*255 / (2)).astype(np.uint8)
        ax.xaxis.set_visible(False)
        ax.yaxis.set_visible(False)
        ax.set_title(age)
        im = ax.imshow(img.reshape((64,64,3)))
```

```
fixed_y_ages = oh_to_class(fixed_y)
_ = view_samples(-1, samples, fixed_y_ages)
```



Vemos que las imágenes generadas que sí parecen caras humanas, si las pone bien, exceptuando las que no parecen

Comienza a programar o generar con IA.